

Classifying-Fashion-MNIST

November 11, 2022

1 Classifying Fashion-MNIST

1.1 Import Resources

```
[1]: import warnings
warnings.filterwarnings('ignore')
```

```
[2]: %matplotlib inline
%config InlineBackend.figure_format = 'retina'

import numpy as np
import matplotlib.pyplot as plt

import tensorflow as tf
import tensorflow_datasets as tfds
tfds.disable_progress_bar()
```

```
[3]: import logging
logger = tf.get_logger()
logger.setLevel(logging.ERROR)
```

```
[4]: print('Using:')
print('\t\u2022 TensorFlow version:', tf.__version__)
print('\t\u2022 tf.keras version:', tf.keras.__version__)
print('\t\u2022 Running on GPU' if tf.test.is_gpu_available() else '\t\u2022 GPU_
    \u2192device not found. Running on CPU')
```

Using:

- TensorFlow version: 2.9.1
- tf.keras version: 2.9.0
- GPU device not found. Running on CPU

1.2 Load the Dataset

```
[6]: train_split = 60
test_val_split = 20
```

```
dataset, dataset_info = tfds.load('fashion_mnist', split=['train[:  
→60%]', 'train[60%:80%]', 'train[80%:]'], as_supervised=True, with_info=True)

training_set, validation_set, test_set = dataset
```

Downloading and preparing dataset Unknown size (download: Unknown size, generated: Unknown size, total: Unknown size) to
C:\Users\LENOVO\tensorflow_datasets\fashion_mnist\3.0.1...
Dataset fashion_mnist downloaded and prepared to
C:\Users\LENOVO\tensorflow_datasets\fashion_mnist\3.0.1. Subsequent calls will reuse this data.

1.3 Explore the Dataset

```
[7]: # Display dataset
dataset
```

```
[7]: [<PrefetchDataset element_spec=(TensorSpec(shape=(28, 28, 1), dtype=tf.uint8, name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>,
      <PrefetchDataset element_spec=(TensorSpec(shape=(28, 28, 1), dtype=tf.uint8, name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>,
      <PrefetchDataset element_spec=(TensorSpec(shape=(28, 28, 1), dtype=tf.uint8, name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>]
```

```
[8]: # Display dataset_info
dataset_info
```

```
[8]: tfds.core.DatasetInfo(
  name='fashion_mnist',
  full_name='fashion_mnist/3.0.1',
  description="""
  Fashion-MNIST is a dataset of Zalando's article images consisting of a
  training set of 60,000 examples and a test set of 10,000 examples. Each example
  is a 28x28 grayscale image, associated with a label from 10 classes.
  """,
  homepage='https://github.com/zalando-research/fashion-mnist',
  data_path='C:\\Users\\LENOVO\\tensorflow_datasets\\fashion_mnist\\3.0.1',
  file_format=tfrecord,
  download_size=29.45 MiB,
  dataset_size=36.42 MiB,
  features=FeaturesDict({
    'image': Image(shape=(28, 28, 1), dtype=tf.uint8),
    'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=10),
  }),
  supervised_keys=('image', 'label'),
```

```

disable_shuffling=False,
splits={
    'test': <SplitInfo num_examples=10000, num_shards=1>,
    'train': <SplitInfo num_examples=60000, num_shards=1>,
},
citation="""@article{DBLP:journals/corr/abs-1708-07747,
  author    = {Han Xiao and
               Kashif Rasul and
               Roland Vollgraf},
  title     = {Fashion-MNIST: a Novel Image Dataset for Benchmarking Machine
Learning
               Algorithms},
  journal   = {CoRR},
  volume    = {abs/1708.07747},
  year      = {2017},
  url       = {http://arxiv.org/abs/1708.07747},
  archivePrefix = {arXiv},
  eprint    = {1708.07747},
  timestamp = {Mon, 13 Aug 2018 16:47:27 +0200},
  biburl    = {https://dblp.org/rec/bib/journals/corr/abs-1708-07747},
  bibsource = {dblp computer science bibliography, https://dblp.org}
}""",
)

```

```

[9]: total_examples = dataset_info.splits['train'].num_examples + dataset_info.
    ↪ splits['test'].num_examples

num_training_examples = (total_examples * train_split) // 100
num_validation_examples = (total_examples * test_val_split) // 100
num_test_examples = num_validation_examples

print('There are {:,} images in the training set'.format(num_training_examples))
print('There are {:,} images in the validation set'.
    ↪ format(num_validation_examples))
print('There are {:,} images in the test set'.format(num_test_examples))

```

There are 42,000 images in the training set
 There are 14,000 images in the validation set
 There are 14,000 images in the test set

The images in this dataset are 28×28 arrays, with pixel values in the range $[0, 255]$. The *labels* are an array of integers, in the range $[0, 9]$. These correspond to the *class* of clothing the image represents:

Label

Class

0

T-shirt/top

1

Trouser

2

Pullover

3

Dress

4

Coat

5

Sandal

6

Shirt

7

Sneaker

8

Bag


9

Ankle boot

Each image is mapped to a single label. Since the *class names* are not included with the dataset, we create them here to use later when plotting the images:

```
[10]: class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',  
                    'Sandal',      'Shirt',   'Sneaker', 'Bag',   'Ankle boot']
```

```
[11]: for image, label in training_set.take(1):  
        image = image.numpy().squeeze()  
        label = label.numpy()  
  
        plt.imshow(image, cmap= plt.cm.binary)  
        plt.colorbar()  
        plt.show()  
  
        print('The label of this image is:', label)  
        print('The class name of this image is:', class_names[label])
```



The label of this image is: 2

The class name of this image is: Pullover

1.4 Create Pipeline

```
[12]: def normalize(image, label):
        image = tf.cast(image, tf.float32)
        image /= 255
        return image, label

batch_size = 64

training_batches = training_set.cache().shuffle(num_training_examples//4).
    ↳ batch(batch_size).map(normalize).prefetch(1)
validation_batches = validation_set.cache().batch(batch_size).map(normalize).
    ↳ prefetch(1)
testing_batches = test_set.cache().batch(batch_size).map(normalize).prefetch(1)
```

1.5 Build the Model

```
[13]: model = tf.keras.Sequential([
        tf.keras.layers.Flatten(input_shape=(28,28,1)),
        tf.keras.layers.Dense(256, activation = 'relu'),
        tf.keras.layers.Dense(128, activation = 'relu'),
        tf.keras.layers.Dense(64, activation = 'relu'),
```

```
tf.keras.layers.Dense(10, activation = 'softmax')
])
```

1.6 Train the Model

```
[14]: model.compile(optimizer='adam',
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

EPOCHS = 30

history = model.fit(training_batches,
                    epochs = EPOCHS,
                    validation_data=validation_batches)
```

```
Epoch 1/30
563/563 [=====] - 4s 6ms/step - loss: 0.5491 -
accuracy: 0.8075 - val_loss: 0.4463 - val_accuracy: 0.8406
Epoch 2/30
563/563 [=====] - 2s 4ms/step - loss: 0.3957 -
accuracy: 0.8559 - val_loss: 0.3826 - val_accuracy: 0.8614
Epoch 3/30
563/563 [=====] - 2s 4ms/step - loss: 0.3508 -
accuracy: 0.8718 - val_loss: 0.3425 - val_accuracy: 0.8752
Epoch 4/30
563/563 [=====] - 2s 4ms/step - loss: 0.3267 -
accuracy: 0.8792 - val_loss: 0.3599 - val_accuracy: 0.8749
Epoch 5/30
563/563 [=====] - 2s 4ms/step - loss: 0.3052 -
accuracy: 0.8871 - val_loss: 0.3421 - val_accuracy: 0.8763
Epoch 6/30
563/563 [=====] - 2s 4ms/step - loss: 0.2877 -
accuracy: 0.8935 - val_loss: 0.3527 - val_accuracy: 0.8728
Epoch 7/30
563/563 [=====] - 2s 4ms/step - loss: 0.2794 -
accuracy: 0.8964 - val_loss: 0.3206 - val_accuracy: 0.8875
Epoch 8/30
563/563 [=====] - 2s 4ms/step - loss: 0.2650 -
accuracy: 0.9005 - val_loss: 0.3600 - val_accuracy: 0.8748
Epoch 9/30
563/563 [=====] - 2s 4ms/step - loss: 0.2507 -
accuracy: 0.9046 - val_loss: 0.3056 - val_accuracy: 0.8884
Epoch 10/30
563/563 [=====] - 2s 4ms/step - loss: 0.2458 -
accuracy: 0.9062 - val_loss: 0.3284 - val_accuracy: 0.8847
Epoch 11/30
563/563 [=====] - 2s 4ms/step - loss: 0.2339 -
accuracy: 0.9112 - val_loss: 0.3377 - val_accuracy: 0.8844
```

Epoch 12/30
563/563 [=====] - 2s 4ms/step - loss: 0.2246 -
accuracy: 0.9138 - val_loss: 0.3506 - val_accuracy: 0.8792
Epoch 13/30
563/563 [=====] - 2s 4ms/step - loss: 0.2183 -
accuracy: 0.9162 - val_loss: 0.3180 - val_accuracy: 0.8888
Epoch 14/30
563/563 [=====] - 2s 4ms/step - loss: 0.2067 -
accuracy: 0.9202 - val_loss: 0.3386 - val_accuracy: 0.8890
Epoch 15/30
563/563 [=====] - 2s 4ms/step - loss: 0.2027 -
accuracy: 0.9233 - val_loss: 0.3259 - val_accuracy: 0.8912
Epoch 16/30
563/563 [=====] - 2s 4ms/step - loss: 0.1941 -
accuracy: 0.9267 - val_loss: 0.3504 - val_accuracy: 0.8842
Epoch 17/30
563/563 [=====] - 2s 4ms/step - loss: 0.1940 -
accuracy: 0.9267 - val_loss: 0.3253 - val_accuracy: 0.8938
Epoch 18/30
563/563 [=====] - 2s 4ms/step - loss: 0.1836 -
accuracy: 0.9293 - val_loss: 0.3766 - val_accuracy: 0.8796
Epoch 19/30
563/563 [=====] - 2s 4ms/step - loss: 0.1723 -
accuracy: 0.9338 - val_loss: 0.3634 - val_accuracy: 0.8912
Epoch 20/30
563/563 [=====] - 2s 4ms/step - loss: 0.1748 -
accuracy: 0.9326 - val_loss: 0.3404 - val_accuracy: 0.8946
Epoch 21/30
563/563 [=====] - 2s 4ms/step - loss: 0.1669 -
accuracy: 0.9351 - val_loss: 0.3740 - val_accuracy: 0.8885
Epoch 22/30
563/563 [=====] - 2s 4ms/step - loss: 0.1639 -
accuracy: 0.9362 - val_loss: 0.3663 - val_accuracy: 0.8919
Epoch 23/30
563/563 [=====] - 2s 4ms/step - loss: 0.1544 -
accuracy: 0.9404 - val_loss: 0.3423 - val_accuracy: 0.8968
Epoch 24/30
563/563 [=====] - 2s 4ms/step - loss: 0.1462 -
accuracy: 0.9428 - val_loss: 0.3737 - val_accuracy: 0.8895
Epoch 25/30
563/563 [=====] - 2s 4ms/step - loss: 0.1436 -
accuracy: 0.9451 - val_loss: 0.3572 - val_accuracy: 0.8932
Epoch 26/30
563/563 [=====] - 2s 4ms/step - loss: 0.1437 -
accuracy: 0.9445 - val_loss: 0.3871 - val_accuracy: 0.8947
Epoch 27/30
563/563 [=====] - 2s 4ms/step - loss: 0.1362 -
accuracy: 0.9465 - val_loss: 0.3790 - val_accuracy: 0.9003

```
Epoch 28/30
563/563 [=====] - 2s 4ms/step - loss: 0.1347 -
accuracy: 0.9472 - val_loss: 0.4164 - val_accuracy: 0.8886
Epoch 29/30
563/563 [=====] - 2s 4ms/step - loss: 0.1233 -
accuracy: 0.9527 - val_loss: 0.3944 - val_accuracy: 0.8943
Epoch 30/30
563/563 [=====] - 2s 4ms/step - loss: 0.1248 -
accuracy: 0.9536 - val_loss: 0.4238 - val_accuracy: 0.8981
```

1.7 Evaluate Loss and Accuracy on the Test Set

```
[15]: loss, accuracy = model.evaluate(testing_batches)

print('\nLoss on the TEST Set: {:.3f}'.format(loss))
print('Accuracy on the TEST Set: {:.3%}'.format(accuracy))

188/188 [=====] - 1s 4ms/step - loss: 0.4134 -
accuracy: 0.8955

Loss on the TEST Set: 0.413
Accuracy on the TEST Set: 89.550%
```

1.8 Loss and Validation Plots


```
[16]: training_accuracy = history.history['accuracy']
validation_accuracy = history.history['val_accuracy']

training_loss = history.history['loss']
validation_loss = history.history['val_loss']

epochs_range=range(EPOCHS)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, training_accuracy, label='Training Accuracy')
plt.plot(epochs_range, validation_accuracy, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, training_loss, label='Training Loss')
plt.plot(epochs_range, validation_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```

1.9 Early Stopping

```
[17]: model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(256, activation = 'relu'),
    tf.keras.layers.Dense(128, activation = 'relu'),
    tf.keras.layers.Dense(64, activation = 'relu'),
    tf.keras.layers.Dense(10, activation = 'softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```

# Stop training when there is no improvement in the validation loss for 5
→consecutive epochs
early_stopping = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5)

history = model.fit(training_batches,
                    epochs = 100,
                    validation_data=validation_batches,
                    callbacks=[early_stopping])

```

```

Epoch 1/100
563/563 [=====] - 3s 4ms/step - loss: 0.5489 -
accuracy: 0.8048 - val_loss: 0.4292 - val_accuracy: 0.8393
Epoch 2/100
563/563 [=====] - 2s 4ms/step - loss: 0.3916 -
accuracy: 0.8575 - val_loss: 0.3855 - val_accuracy: 0.8584
Epoch 3/100
563/563 [=====] - 2s 4ms/step - loss: 0.3534 -
accuracy: 0.8706 - val_loss: 0.3741 - val_accuracy: 0.8683
Epoch 4/100
563/563 [=====] - 3s 5ms/step - loss: 0.3265 -
accuracy: 0.8806 - val_loss: 0.3523 - val_accuracy: 0.8729
Epoch 5/100
563/563 [=====] - 2s 4ms/step - loss: 0.3066 -
accuracy: 0.8859 - val_loss: 0.3425 - val_accuracy: 0.8780
Epoch 6/100
563/563 [=====] - 2s 4ms/step - loss: 0.2892 -
accuracy: 0.8907 - val_loss: 0.3599 - val_accuracy: 0.8725
Epoch 7/100
563/563 [=====] - 2s 4ms/step - loss: 0.2758 -
accuracy: 0.8975 - val_loss: 0.3409 - val_accuracy: 0.8748
Epoch 8/100
563/563 [=====] - 2s 4ms/step - loss: 0.2653 -
accuracy: 0.9003 - val_loss: 0.3200 - val_accuracy: 0.8865
Epoch 9/100
563/563 [=====] - 2s 4ms/step - loss: 0.2540 -
accuracy: 0.9040 - val_loss: 0.3253 - val_accuracy: 0.8867
Epoch 10/100
563/563 [=====] - 2s 4ms/step - loss: 0.2475 -
accuracy: 0.9053 - val_loss: 0.3540 - val_accuracy: 0.8785
Epoch 11/100
563/563 [=====] - 2s 4ms/step - loss: 0.2350 -
accuracy: 0.9115 - val_loss: 0.3224 - val_accuracy: 0.8870
Epoch 12/100
563/563 [=====] - 3s 4ms/step - loss: 0.2275 -
accuracy: 0.9150 - val_loss: 0.3345 - val_accuracy: 0.8867
Epoch 13/100
563/563 [=====] - 2s 4ms/step - loss: 0.2202 -

```

accuracy: 0.9184 - val_loss: 0.3302 - val_accuracy: 0.8876


```
[18]: training_accuracy = history.history['accuracy']
      validation_accuracy = history.history['val_accuracy']

      training_loss = history.history['loss']
      validation_loss = history.history['val_loss']

      epochs_range=range(len(training_accuracy))

      plt.figure(figsize=(8, 8))
      plt.subplot(1, 2, 1)
      plt.plot(epochs_range, training_accuracy, label='Training Accuracy')
      plt.plot(epochs_range, validation_accuracy, label='Validation Accuracy')
      plt.legend(loc='lower right')
      plt.title('Training and Validation Accuracy')

      plt.subplot(1, 2, 2)
      plt.plot(epochs_range, training_loss, label='Training Loss')
      plt.plot(epochs_range, validation_loss, label='Validation Loss')
      plt.legend(loc='upper right')
      plt.title('Training and Validation Loss')
      plt.show()
```



1.10 Dropout

```
[19]: model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(256, activation = 'relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(128, activation = 'relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(64, activation = 'relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation = 'softmax')
])
```

```

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(training_batches,
                    epochs = 30,
                    validation_data=validation_batches)

```

```

Epoch 1/30
563/563 [=====] - 3s 5ms/step - loss: 0.6822 -
accuracy: 0.7566 - val_loss: 0.4395 - val_accuracy: 0.8422
Epoch 2/30
563/563 [=====] - 3s 5ms/step - loss: 0.4745 -
accuracy: 0.8324 - val_loss: 0.4032 - val_accuracy: 0.8525
Epoch 3/30
563/563 [=====] - 3s 5ms/step - loss: 0.4240 -
accuracy: 0.8461 - val_loss: 0.3955 - val_accuracy: 0.8533
Epoch 4/30
563/563 [=====] - 3s 5ms/step - loss: 0.4002 -
accuracy: 0.8557 - val_loss: 0.3587 - val_accuracy: 0.8688
Epoch 5/30
563/563 [=====] - 3s 5ms/step - loss: 0.3741 -
accuracy: 0.8650 - val_loss: 0.3556 - val_accuracy: 0.8707
Epoch 6/30
563/563 [=====] - 3s 5ms/step - loss: 0.3618 -
accuracy: 0.8681 - val_loss: 0.3711 - val_accuracy: 0.8612
Epoch 7/30
563/563 [=====] - 3s 5ms/step - loss: 0.3498 -
accuracy: 0.8711 - val_loss: 0.3414 - val_accuracy: 0.8777
Epoch 8/30
563/563 [=====] - 3s 5ms/step - loss: 0.3386 -
accuracy: 0.8760 - val_loss: 0.3531 - val_accuracy: 0.8701
Epoch 9/30
563/563 [=====] - 3s 5ms/step - loss: 0.3272 -
accuracy: 0.8809 - val_loss: 0.3199 - val_accuracy: 0.8861
Epoch 10/30
563/563 [=====] - 3s 5ms/step - loss: 0.3223 -
accuracy: 0.8819 - val_loss: 0.3417 - val_accuracy: 0.8816
Epoch 11/30
563/563 [=====] - 3s 4ms/step - loss: 0.3092 -
accuracy: 0.8852 - val_loss: 0.3225 - val_accuracy: 0.8848
Epoch 12/30
563/563 [=====] - 3s 5ms/step - loss: 0.3016 -
accuracy: 0.8890 - val_loss: 0.3295 - val_accuracy: 0.8830
Epoch 13/30
563/563 [=====] - 3s 5ms/step - loss: 0.3023 -
accuracy: 0.8898 - val_loss: 0.3539 - val_accuracy: 0.8767

```

Epoch 14/30
563/563 [=====] - 3s 5ms/step - loss: 0.2896 - accuracy: 0.8936 - val_loss: 0.3262 - val_accuracy: 0.8863

Epoch 15/30
563/563 [=====] - 3s 5ms/step - loss: 0.2865 - accuracy: 0.8935 - val_loss: 0.3336 - val_accuracy: 0.8867

Epoch 16/30
563/563 [=====] - 3s 5ms/step - loss: 0.2818 - accuracy: 0.8956 - val_loss: 0.3223 - val_accuracy: 0.8871

Epoch 17/30
563/563 [=====] - 3s 5ms/step - loss: 0.2748 - accuracy: 0.8971 - val_loss: 0.3252 - val_accuracy: 0.8891

Epoch 18/30
563/563 [=====] - 3s 5ms/step - loss: 0.2733 - accuracy: 0.8979 - val_loss: 0.3103 - val_accuracy: 0.8906

Epoch 19/30
563/563 [=====] - 3s 5ms/step - loss: 0.2693 - accuracy: 0.9004 - val_loss: 0.3266 - val_accuracy: 0.8879

Epoch 20/30
563/563 [=====] - 3s 5ms/step - loss: 0.2650 - accuracy: 0.9017 - val_loss: 0.3255 - val_accuracy: 0.8920

Epoch 21/30
563/563 [=====] - 3s 5ms/step - loss: 0.2589 - accuracy: 0.9038 - val_loss: 0.3233 - val_accuracy: 0.8829

Epoch 22/30
563/563 [=====] - 3s 5ms/step - loss: 0.2539 - accuracy: 0.9051 - val_loss: 0.3131 - val_accuracy: 0.8922

Epoch 23/30
563/563 [=====] - 3s 5ms/step - loss: 0.2530 - accuracy: 0.9075 - val_loss: 0.3279 - val_accuracy: 0.8831

Epoch 24/30
563/563 [=====] - 3s 5ms/step - loss: 0.2437 - accuracy: 0.9073 - val_loss: 0.3240 - val_accuracy: 0.8917

Epoch 25/30
563/563 [=====] - 3s 5ms/step - loss: 0.2441 - accuracy: 0.9089 - val_loss: 0.3177 - val_accuracy: 0.8928

Epoch 26/30
563/563 [=====] - 3s 5ms/step - loss: 0.2455 - accuracy: 0.9081 - val_loss: 0.3108 - val_accuracy: 0.8942

Epoch 27/30
563/563 [=====] - 3s 5ms/step - loss: 0.2381 - accuracy: 0.9111 - val_loss: 0.3187 - val_accuracy: 0.8917

Epoch 28/30
563/563 [=====] - 3s 5ms/step - loss: 0.2347 - accuracy: 0.9120 - val_loss: 0.3259 - val_accuracy: 0.8918

Epoch 29/30
563/563 [=====] - 3s 5ms/step - loss: 0.2353 - accuracy: 0.9109 - val_loss: 0.3122 - val_accuracy: 0.8937

Epoch 30/30

563/563 [=====] - 3s 5ms/step - loss: 0.2272 - accuracy: 0.9147 - val_loss: 0.3146 - val_accuracy: 0.8900


```
[20]: training_accuracy = history.history['accuracy']
      validation_accuracy = history.history['val_accuracy']

      training_loss = history.history['loss']
      validation_loss = history.history['val_loss']

      epochs_range=range(len(training_accuracy))

      plt.figure(figsize=(8, 8))
      plt.subplot(1, 2, 1)
      plt.plot(epochs_range, training_accuracy, label='Training Accuracy')
      plt.plot(epochs_range, validation_accuracy, label='Validation Accuracy')
      plt.legend(loc='lower right')
      plt.title('Training and Validation Accuracy')

      plt.subplot(1, 2, 2)
      plt.plot(epochs_range, training_loss, label='Training Loss')
      plt.plot(epochs_range, validation_loss, label='Validation Loss')
      plt.legend(loc='upper right')
      plt.title('Training and Validation Loss')
      plt.show()
```



1.11 Inference

```
[21]: for image_batch, label_batch in testing_batches.take(1):
        ps = model.predict(image_batch)
        images = image_batch.numpy().squeeze()
        labels = label_batch.numpy()

plt.figure(figsize=(10,15))

for n in range(30):
    plt.subplot(6,5,n+1)
    plt.imshow(images[n], cmap = plt.cm.binary)
```



```
color = 'green' if np.argmax(ps[n]) == labels[n] else 'red'  
plt.title(class_names[np.argmax(ps[n])], color=color)  
plt.axis('off')
```

2/2 [=====] - 0s 3ms/step

Sneaker



Shirt




Sandal




Sneaker



Pullover



Pullover



T-shirt/top



Dress



Sandal



Coat



Bag



Bag



Dress



Sandal



T-shirt/top



Shirt



Bag



Dress



Sneaker



Ankle boot



Coat



Sneaker



Sneaker



Ankle boot



T-shirt/top



Coat



T-shirt/top



Sneaker



T-shirt/top



Ankle boot



[]: