

▼ import libraries

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
import torch.backends.cudnn as cudnn
import torchvision
from torchvision import datasets, models, transforms
import matplotlib.pyplot as plt
import numpy as np
import os, sys
import time
import copy
from glob import glob
import imageio
```

▼ Download Intel Image Dataset using Kaggle api

```
!pip install kaggle
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Requirement already satisfied: kaggle in /usr/local/lib/python3.8/dist-packages (1.5.12)
Requirement already satisfied: certifi in /usr/local/lib/python3.8/dist-packages (from kaggle) (2022.9.24)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.8/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.8/dist-packages (from kaggle) (2.28.1)
Requirement already satisfied: tqdm in /usr/local/lib/python3.8/dist-packages (from kaggle) (4.64.1)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.8/dist-packages (from kaggle) (5.0.2)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.8/dist-packages (from kaggle) (1.26.12)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.8/dist-packages (from kaggle) (1.16.0)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.8/dist-packages (from kaggle) (1.3)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from kaggle) (3.4)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from kaggle) (3.7.4)
```

```
!mkdir ~/.kaggle
!cp kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json
```

```
!kaggle datasets download -d puneet6060/intel-image-classification
```

```
Downloading intel-image-classification.zip to /content
100% 346M/346M [00:16<00:00, 20.9MB/s]
100% 346M/346M [00:16<00:00, 20.9MB/s]
```

1s completed at 9:32 PM



▼ Data directory set

```
!unzip intel-image-classification.zip
```

Streaming output truncated to the last 5000 lines.

```
inflating: seg_train/seg_train/mountain/7506.jpg
inflating: seg_train/seg_train/mountain/7537.jpg
inflating: seg_train/seg_train/mountain/7539.jpg
inflating: seg_train/seg_train/mountain/7551.jpg
inflating: seg_train/seg_train/mountain/7560.jpg
inflating: seg_train/seg_train/mountain/7565.jpg
inflating: seg_train/seg_train/mountain/7578.jpg
inflating: seg_train/seg_train/mountain/7581.jpg
inflating: seg_train/seg_train/mountain/7586.jpg
inflating: seg_train/seg_train/mountain/7647.jpg
inflating: seg_train/seg_train/mountain/7652.jpg
inflating: seg_train/seg_train/mountain/7654.jpg
inflating: seg_train/seg_train/mountain/7662.jpg
inflating: seg_train/seg_train/mountain/767.jpg
inflating: seg_train/seg_train/mountain/7672.jpg
inflating: seg_train/seg_train/mountain/7679.jpg
inflating: seg_train/seg_train/mountain/7681.jpg
inflating: seg_train/seg_train/mountain/7693.jpg
inflating: seg_train/seg_train/mountain/7695.jpg
inflating: seg_train/seg_train/mountain/7698.jpg
inflating: seg_train/seg_train/mountain/7700.jpg
inflating: seg_train/seg_train/mountain/771.jpg
inflating: seg_train/seg_train/mountain/7715.jpg
inflating: seg_train/seg_train/mountain/7744.jpg
inflating: seg_train/seg_train/mountain/7745.jpg
inflating: seg_train/seg_train/mountain/7751.jpg
inflating: seg_train/seg_train/mountain/7763.jpg
inflating: seg_train/seg_train/mountain/7771.jpg
inflating: seg_train/seg_train/mountain/7780.jpg
inflating: seg_train/seg_train/mountain/7787.jpg
inflating: seg_train/seg_train/mountain/7788.jpg
inflating: seg_train/seg_train/mountain/7813.jpg
inflating: seg_train/seg_train/mountain/7816.jpg
inflating: seg_train/seg_train/mountain/7819.jpg
inflating: seg_train/seg_train/mountain/7820.jpg
inflating: seg_train/seg_train/mountain/7823.jpg
inflating: seg_train/seg_train/mountain/7836.jpg
inflating: seg_train/seg_train/mountain/784.jpg
inflating: seg_train/seg_train/mountain/7841.jpg
inflating: seg_train/seg_train/mountain/7842.jpg
inflating: seg_train/seg_train/mountain/7845.jpg
inflating: seg_train/seg_train/mountain/7849.jpg
inflating: seg_train/seg_train/mountain/7851.jpg
inflating: seg_train/seg_train/mountain/7865.jpg
inflating: seg_train/seg_train/mountain/7875.jpg
```

```
inflating: seg_train/seg_train/mountain/7881.jpg
inflating: seg_train/seg_train/mountain/7885.jpg
inflating: seg_train/seg_train/mountain/790.jpg
inflating: seg_train/seg_train/mountain/7908.jpg
inflating: seg_train/seg_train/mountain/7909.jpg
inflating: seg_train/seg_train/mountain/7912.jpg
inflating: seg_train/seg_train/mountain/7922.jpg
inflating: seg_train/seg_train/mountain/7928.jpg
inflating: seg_train/seg_train/mountain/7942.jpg
inflating: seg_train/seg_train/mountain/7946.jpg
inflating: seg_train/seg_train/mountain/7960.jpg
inflating: seg_train/seg_train/mountain/7973.jpg
```

```
train = '/content/seg_train/seg_train'
test = '/content/seg_test/seg_test'
```

Data Augmentation

Note: normalize mean and std are standardized for ImageNet

```
data_transforms = {
    train: transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(0.5),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    test: transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
}
```

Load data

```
image_datasets = {x: datasets.ImageFolder(os.path.join(x), data_transforms[x]) for x in [train, test]}
dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=32, shuffle=True) for x in [train, test]}
dataset_sizes = {x: len(image_datasets[x]) for x in [train, test]}
```

```

class_names = image_datasets[train].classes

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554: UserWarning
  warnings.warn(_create_warning_msg(

resnet = models.resnet152(pretrained=True)

/usr/local/lib/python3.8/dist-packages/torchvision/models/_utils.py:208: UserWarning
  warnings.warn(
/usr/local/lib/python3.8/dist-packages/torchvision/models/_utils.py:223: UserWarning
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet152-394f9c45.pth" to /root/.
100%                230M/230M [00:00<00:00, 272MB/s]

```

Plot intel image dataset

```

def imshow(inp, title=None):
    """Imshow for Tensor."""
    inp = inp.numpy().transpose((1, 2, 0))
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])
    inp = std * inp + mean
    inp = np.clip(inp, 0, 1)
    plt.imshow(inp)
    if title is not None:
        plt.title(title)
    plt.pause(0.001) # pause a bit so that plots are updated

```

```

# Get a batch of training data
inputs, classes = next(iter(dataloader[train]))

```

```

# Make a grid from batch
out = torchvision.utils.make_grid(inputs)

```

```

imshow(out, title=[class_names[x] for x in classes])

```



```
print(image_datasets[train])

Dataset ImageFolder
  Number of datapoints: 14034
  Root location: /content/seg_train/seg_train
  StandardTransform
Transform: Compose(
  RandomResizedCrop(size=(224, 224), scale=(0.08, 1.0), ratio=(0.75, 1.0))
  RandomHorizontalFlip(p=0.5)
  ToTensor()
  Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
)
```

Freeze CNN part of ResNet152 Model

```
for param in resnet.parameters():
    param.requires_grad = False
```

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

Set fully connected sequential

```
resnet.fc = nn.Sequential(nn.Linear(2048, 1024),
                          nn.ReLU(),
                          nn.Dropout(0.25),
                          nn.Linear(1024, 256),
                          nn.ReLU(),
                          nn.Dropout(0.25),
                          nn.Linear(256, 6),
                          nn.LogSoftmax(dim=1))
```

```
resnet = resnet.to(device)
```

Loss function/ Optimizer

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(resnet.parameters(), lr=1e-4)
exp_lr_scheduler = lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)
```

Train Model Function

```
from datetime import datetime

def train_model(model, criterion, optimizer, train_loader, test_loader, epochs):
    train_losses = np.zeros(epochs)
    test_losses = np.zeros(epochs)

    for it in range(epochs):
        if it == 10 :
            for param in model.parameters():
                param.requires_grad = True
        t0 = datetime.now()
        train_loss = []
        for inputs, targets in dataloader[train]:
            inputs, targets = inputs.to(device), targets.to(device)
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, targets)
            loss.backward()
            optimizer.step()
            train_loss.append(loss.item())

        train_loss = np.mean(train_loss)

        test_loss = []
        for inputs, targets in dataloader[test]:
            inputs, targets = inputs.to(device), targets.to(device)
            outputs = model(inputs)
            loss = criterion(outputs, targets)
            test_loss.append(loss.item())

        test_loss = np.mean(test_loss)

        train_losses[it] = train_loss
        test_losses[it] = test_loss

        dt = datetime.now() - t0

        print(f'Epoch {it+1}/{epochs}, Train_Loss: {train_loss:.4f}, \Test_Loss: {test_loss:.4f}')

    return train_losses, test_losses

train_losses, test_losses = train_model(
    resnet,
    criterion.
```


```
optimizer,  
dataloader[train],  
dataloader[test],  
epochs=40)
```

```
Epoch 1/40, Train_Loss: 0.6348, \Test_Loss: 0.3671, Duration: 0:02:18.012346  
Epoch 2/40, Train_Loss: 0.4187, \Test_Loss: 0.3239, Duration: 0:02:12.348367  
Epoch 3/40, Train_Loss: 0.4116, \Test_Loss: 0.3183, Duration: 0:02:12.303492  
Epoch 4/40, Train_Loss: 0.3940, \Test_Loss: 0.3011, Duration: 0:02:11.797530  
Epoch 5/40, Train_Loss: 0.3809, \Test_Loss: 0.2869, Duration: 0:02:12.182279  
Epoch 6/40, Train_Loss: 0.3760, \Test_Loss: 0.3113, Duration: 0:02:12.008638  
Epoch 7/40, Train_Loss: 0.3771, \Test_Loss: 0.2661, Duration: 0:02:12.230927  
Epoch 8/40, Train_Loss: 0.3743, \Test_Loss: 0.2872, Duration: 0:02:12.286634  
Epoch 9/40, Train_Loss: 0.3530, \Test_Loss: 0.2662, Duration: 0:02:13.266736  
Epoch 10/40, Train_Loss: 0.3539, \Test_Loss: 0.2692, Duration: 0:02:12.489712  
Epoch 11/40, Train_Loss: 0.4025, \Test_Loss: 0.2725, Duration: 0:06:02.206394  
Epoch 12/40, Train_Loss: 0.3366, \Test_Loss: 0.2625, Duration: 0:06:00.994119  
Epoch 13/40, Train_Loss: 0.3119, \Test_Loss: 0.2546, Duration: 0:06:00.693115  
Epoch 14/40, Train_Loss: 0.2928, \Test_Loss: 0.2227, Duration: 0:06:01.171837  
Epoch 15/40, Train_Loss: 0.2811, \Test_Loss: 0.2380, Duration: 0:06:00.407107  
Epoch 16/40, Train_Loss: 0.2797, \Test_Loss: 0.2192, Duration: 0:06:00.845971  
Epoch 17/40, Train_Loss: 0.2697, \Test_Loss: 0.2277, Duration: 0:06:00.459688  
Epoch 18/40, Train_Loss: 0.2573, \Test_Loss: 0.2326, Duration: 0:06:00.712149  
Epoch 19/40, Train_Loss: 0.2453, \Test_Loss: 0.2225, Duration: 0:06:00.095190  
Epoch 20/40, Train_Loss: 0.2413, \Test_Loss: 0.2334, Duration: 0:06:00.221754  
Epoch 21/40, Train_Loss: 0.2371, \Test_Loss: 0.2379, Duration: 0:05:59.689844  
Epoch 22/40, Train_Loss: 0.2397, \Test_Loss: 0.2109, Duration: 0:05:59.977015  
Epoch 23/40, Train_Loss: 0.2248, \Test_Loss: 0.2046, Duration: 0:05:59.677062  
Epoch 24/40, Train_Loss: 0.2104, \Test_Loss: 0.2154, Duration: 0:05:59.551734  
Epoch 25/40, Train_Loss: 0.2113, \Test_Loss: 0.2479, Duration: 0:05:59.383888  
Epoch 26/40, Train_Loss: 0.2074, \Test_Loss: 0.2388, Duration: 0:05:59.620912  
Epoch 27/40, Train_Loss: 0.2065, \Test_Loss: 0.2240, Duration: 0:05:59.554906  
Epoch 28/40, Train_Loss: 0.2023, \Test_Loss: 0.2087, Duration: 0:05:57.515687  
Epoch 29/40, Train_Loss: 0.1871, \Test_Loss: 0.2478, Duration: 0:05:55.495677  
Epoch 30/40, Train_Loss: 0.1866, \Test_Loss: 0.2309, Duration: 0:05:55.362354  
Epoch 31/40, Train_Loss: 0.1797, \Test_Loss: 0.2314, Duration: 0:05:55.285163  
Epoch 32/40, Train_Loss: 0.1787, \Test_Loss: 0.2256, Duration: 0:05:54.940069  
Epoch 33/40, Train_Loss: 0.1660, \Test_Loss: 0.2410, Duration: 0:05:56.998170  
Epoch 34/40, Train_Loss: 0.1751, \Test_Loss: 0.2341, Duration: 0:05:55.879126  
Epoch 35/40, Train_Loss: 0.1724, \Test_Loss: 0.2451, Duration: 0:05:55.100177  
Epoch 36/40, Train_Loss: 0.1680, \Test_Loss: 0.2369, Duration: 0:05:54.989463  
Epoch 37/40, Train_Loss: 0.1628, \Test_Loss: 0.2321, Duration: 0:05:55.205754  
Epoch 38/40, Train_Loss: 0.1506, \Test_Loss: 0.2326, Duration: 0:05:54.941748  
Epoch 39/40, Train_Loss: 0.1547, \Test_Loss: 0.2317, Duration: 0:05:54.979370  
Epoch 40/40, Train_Loss: 0.1539, \Test_Loss: 0.2552, Duration: 0:05:54.774343
```

Plot Train/Test loss

```
plt.plot(train_losses, label='train loss')  
plt.plot(test_losses, label='test loss')  
plt.legend()
```

```
plt.legend()
plt.show()
```



Accuracy on Test data

```
from tqdm.autonotebook import tqdm
```

```
train_losses=[]
test_losses=[]
def accuracy(loader, model):
    num_corrects = 0
    num_samples = 0
    model.eval()
    loop = tqdm(loader)
    with torch.no_grad():
        for x, y in loop:
            x = x.to(device)
            y = y.to(device)
            scores = model(x)
            test_losses.append(scores.data)
            _, prediction = scores.max(1)
            num_corrects += (prediction == y).sum()
            num_samples += prediction.size(0)
            acc = (num_corrects/num_samples) * 100
            loop.set_postfix(acc=acc.item())
    print(f'Got {num_corrects}/{num_samples} with accuracy {acc:.4f}')
```

```
accuracy(dataloader[test], resnet)
```

100%

94/94 [00:21<00:00, 4.93it/s, acc=93.1]

Got 2793/3000 with accuracy 93.1000

Train dataset with Wide Residual Networks (WRNs)

```
wrn = models.wide_resnet101_2(pretrained=True)
```

```
/usr/local/lib/python3.8/dist-packages/torchvision/models/_utils.py:223: UserWarning
warnings.warn(msg)
```

```
Downloading: "https://download.pytorch.org/models/wide\_resnet101\_2-32ee1156.pth" to .
```

```
100%
```

```
243M/243M [00:01<00:00, 276MB/s]
```

```
for param in wrn.parameters():
    param.requires_grad = False
```

```
wrn.fc = nn.Sequential(nn.Linear(2048, 1024),
                       nn.ReLU(),
                       nn.Dropout(0.5),
                       nn.Linear(1024, 6),
                       nn.LogSoftmax(dim=1))
```

```
wrn = wrn.to(device)
```

```
data_transforms = {
    train: transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(0.5),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    test: transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
}
```

```
image_datasets = {x: datasets.ImageFolder(os.path.join(x), data_transforms[x]) for x in [t
```

```
dataloader = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=32, shuffle=True
```

```
dataset_sizes = {x: len(image_datasets[x]) for x in [train,test]}
```

```
class_names = image_datasets[train].classes
```

```
/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554: UserWarning:
```

```
/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554: UserWarning:
  warnings.warn(_create_warning_msg(
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(wrn.parameters(), lr=1e-3, betas=(0.9, 0.9))

import torchsummary as summary

def train_model(model, criterion, optimizer, train_loader, test_loader, epochs):
    train_losses = np.zeros(epochs)
    test_losses = np.zeros(epochs)

    for it in range(epochs):
        t0 = datetime.now()
        train_loss = []
        for inputs, targets in dataloader[train]:
            inputs, targets = inputs.to(device), targets.to(device)
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, targets)
            loss.backward()
            optimizer.step()
            train_loss.append(loss.item())

        train_loss = np.mean(train_loss)

        test_loss = []
        for inputs, targets in dataloader[test]:
            inputs, targets = inputs.to(device), targets.to(device)
            outputs = model(inputs)
            loss = criterion(outputs, targets)
            test_loss.append(loss.item())

        test_loss = np.mean(test_loss)

        train_losses[it] = train_loss
        test_losses[it] = test_loss

        dt = datetime.now() - t0

        print(f'Epoch {it+1}/{epochs}, Train_Loss: {train_loss:.4f}, \Test_Loss: {test_loss:.4f}')


    return train_losses, test_losses

train_losses, test_losses = train_model(
    wrn,
    criterion,
    optimizer,
    train_loader,
    test_loader,
    epochs)
```

```
dataloader[train],  
dataloader[test],  
epochs=30)
```

```
Epoch 1/30, Train_Loss: 0.5889, \Test_Loss: 0.4229, Duration: 0:02:47.483557  
Epoch 2/30, Train_Loss: 0.5284, \Test_Loss: 0.3889, Duration: 0:02:53.000433  
Epoch 3/30, Train_Loss: 0.5128, \Test_Loss: 0.4088, Duration: 0:02:55.257686  
Epoch 4/30, Train_Loss: 0.5026, \Test_Loss: 0.3822, Duration: 0:02:56.183388  
Epoch 5/30, Train_Loss: 0.4924, \Test_Loss: 0.3778, Duration: 0:02:55.113532  
Epoch 6/30, Train_Loss: 0.4870, \Test_Loss: 0.3805, Duration: 0:02:54.972858  
Epoch 7/30, Train_Loss: 0.4976, \Test_Loss: 0.3897, Duration: 0:02:55.538814  
Epoch 8/30, Train_Loss: 0.4827, \Test_Loss: 0.3687, Duration: 0:02:55.844843  
Epoch 9/30, Train_Loss: 0.4908, \Test_Loss: 0.3762, Duration: 0:02:56.095436  
Epoch 10/30, Train_Loss: 0.5021, \Test_Loss: 0.3859, Duration: 0:02:55.548505  
Epoch 11/30, Train_Loss: 0.4970, \Test_Loss: 0.3798, Duration: 0:02:55.285072  
Epoch 12/30, Train_Loss: 0.4954, \Test_Loss: 0.3874, Duration: 0:02:55.819190  
Epoch 13/30, Train_Loss: 0.4992, \Test_Loss: 0.3777, Duration: 0:02:55.143227  
Epoch 14/30, Train_Loss: 0.5024, \Test_Loss: 0.3818, Duration: 0:02:55.304871  
Epoch 15/30, Train_Loss: 0.5091, \Test_Loss: 0.3609, Duration: 0:02:55.338294  
Epoch 16/30, Train_Loss: 0.5090, \Test_Loss: 0.3478, Duration: 0:02:56.312939  
Epoch 17/30, Train_Loss: 0.4913, \Test_Loss: 0.3913, Duration: 0:02:55.157252  
Epoch 18/30, Train_Loss: 0.5038, \Test_Loss: 0.3716, Duration: 0:02:55.323865  
Epoch 19/30, Train_Loss: 0.5054, \Test_Loss: 0.3706, Duration: 0:02:55.515696  
Epoch 20/30, Train_Loss: 0.5124, \Test_Loss: 0.3981, Duration: 0:02:55.673086  
Epoch 21/30, Train_Loss: 0.5091, \Test_Loss: 0.3823, Duration: 0:02:55.518970  
Epoch 22/30, Train_Loss: 0.5055, \Test_Loss: 0.4230, Duration: 0:02:55.496015  
Epoch 23/30, Train_Loss: 0.5018, \Test_Loss: 0.4120, Duration: 0:02:55.743689  
Epoch 24/30, Train_Loss: 0.5240, \Test_Loss: 0.4433, Duration: 0:02:55.354273  
Epoch 25/30, Train_Loss: 0.5006, \Test_Loss: 0.4176, Duration: 0:02:55.864043  
Epoch 26/30, Train_Loss: 0.5179, \Test_Loss: 0.3961, Duration: 0:02:56.052837  
Epoch 27/30, Train_Loss: 0.5091, \Test_Loss: 0.3867, Duration: 0:02:55.890951  
Epoch 28/30, Train_Loss: 0.5107, \Test_Loss: 0.3880, Duration: 0:02:55.322546  
Epoch 29/30, Train_Loss: 0.5188, \Test_Loss: 0.4006, Duration: 0:02:55.645223  
Epoch 30/30, Train_Loss: 0.4990, \Test_Loss: 0.3869, Duration: 0:02:55.779137
```

```
plt.plot(train_losses, label='train loss')  
plt.plot(test_losses, label='test loss')  
plt.legend()  
plt.show()
```



```
n_correct = 0
n_total = 0
for inputs, targets in dataloader[train]:
    inputs, targets = inputs.to(device), targets.to(device)
    outputs = wrn(inputs)
    _, predictions = torch.max(outputs, -1)
    n_correct += (predictions == targets).sum().item()
    n_total += targets.shape[0]
train_acc = n_correct/n_total

n_correct = 0
n_total = 0
for inputs, targets in dataloader[test]:
    inputs, targets = inputs.to(device), targets.to(device)
    outputs = wrn(inputs)
    _, predictions = torch.max(outputs, -1)
    n_correct += (predictions == targets).sum().item()
    n_total += targets.shape[0]
test_acc = n_correct/n_total
print(f'Train acc: {train_acc:.4f}, Test acc: {test_acc:.4f}')
```

```
/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554: UserWarning:
  warnings.warn(_create_warning_msg(
Train acc: 0.8312, Test acc: 0.8653
```

```
train_losses=[]
test_losses=[]
def accuracy(loader, model):
    num_corrects = 0
    num_samples = 0
    model.eval()
    loop = tqdm(loader)
    with torch.no_grad():
        for x, y in loop:
            x = x.to(device)
            y = y.to(device)
            scores = model(x)
            test_losses.append(scores.data)
            _, prediction = scores.max(1)
            num_corrects += (prediction == y).sum()
            num_samples += prediction.size(0)
            acc = (num_corrects/num_samples) * 100
            loop.set_postfix(acc=acc.item())
        print(f'Got {num_corrects}/{num_samples} with accuracy {acc:.4f}')
```

```
accuracy(dataloader[test], wrn)
```

100%

94/94 [00:29<00:00, 3.54it/s, acc=89.5]

```
/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554: UserWarning:
  warnings.warn(_create_warning_msg(
Got 2685/3000 with accuracy 89.5000
```

Train intel-image using Inception-ResNet

```
model = models.inception_v3(pretrained=True)
```

```
def get_model():
    model = models.inception_v3(pretrained=True)
    for param in model.parameters():
        param.requires_grad = False #Freezing all the layers and changing only the below 1
    model.avgpool = nn.AdaptiveAvgPool2d(output_size=(1,1))
    model.fc = nn.Sequential(nn.Flatten(),
                             nn.Linear(2048,1024),
                             nn.ReLU(),
                             nn.Dropout(0.3),
                             nn.Linear(1024,256),
                             nn.ReLU(),
                             nn.Dropout(0.3),
                             nn.Linear(256,6),
                             nn.LogSoftmax(dim=1))

    model.aux_logits = False
    loss_fn = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=1e-4, betas=(0.9, 0.9))
    return model.to(device), loss_fn, optimizer
```

```
!pip install torchsummary
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Requirement already satisfied: torchsummary in /usr/local/lib/python3.8/dist-package
```

```
from torchsummary import summary
```

```
input_shape = (3,300,300)
summary(model.to(device), input_shape)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 149, 149]	864
BatchNorm2d-2	[-1, 32, 149, 149]	64
BasicConv2d-3	[-1, 32, 149, 149]	0
Conv2d-4	[-1, 32, 147, 147]	9,216
BatchNorm2d-5	[-1, 32, 147, 147]	64
BasicConv2d-6	[-1, 32, 147, 147]	0
Conv2d-7	[-1, 64, 147, 147]	18,432
BatchNorm2d-8	[-1, 64, 147, 147]	128
BasicConv2d-9	[-1, 64, 147, 147]	0
MaxPool2d-10	[-1, 64, 73, 73]	0
Conv2d-11	[-1, 80, 73, 73]	5,120
BatchNorm2d-12	[-1, 80, 73, 73]	160
BasicConv2d-13	[-1, 80, 73, 73]	0
Conv2d-14	[-1, 192, 71, 71]	138,240
BatchNorm2d-15	[-1, 192, 71, 71]	384
BasicConv2d-16	[-1, 192, 71, 71]	0
MaxPool2d-17	[-1, 192, 35, 35]	0
Conv2d-18	[-1, 64, 35, 35]	12,288
BatchNorm2d-19	[-1, 64, 35, 35]	128
BasicConv2d-20	[-1, 64, 35, 35]	0
Conv2d-21	[-1, 48, 35, 35]	9,216
BatchNorm2d-22	[-1, 48, 35, 35]	96
BasicConv2d-23	[-1, 48, 35, 35]	0
Conv2d-24	[-1, 64, 35, 35]	76,800
BatchNorm2d-25	[-1, 64, 35, 35]	128
BasicConv2d-26	[-1, 64, 35, 35]	0
Conv2d-27	[-1, 64, 35, 35]	12,288
BatchNorm2d-28	[-1, 64, 35, 35]	128
BasicConv2d-29	[-1, 64, 35, 35]	0
Conv2d-30	[-1, 96, 35, 35]	55,296
BatchNorm2d-31	[-1, 96, 35, 35]	192
BasicConv2d-32	[-1, 96, 35, 35]	0
Conv2d-33	[-1, 96, 35, 35]	82,944
BatchNorm2d-34	[-1, 96, 35, 35]	192
BasicConv2d-35	[-1, 96, 35, 35]	0
Conv2d-36	[-1, 32, 35, 35]	6,144
BatchNorm2d-37	[-1, 32, 35, 35]	64
BasicConv2d-38	[-1, 32, 35, 35]	0
InceptionA-39	[-1, 256, 35, 35]	0
Conv2d-40	[-1, 64, 35, 35]	16,384
BatchNorm2d-41	[-1, 64, 35, 35]	128
BasicConv2d-42	[-1, 64, 35, 35]	0
Conv2d-43	[-1, 48, 35, 35]	12,288
BatchNorm2d-44	[-1, 48, 35, 35]	96
BasicConv2d-45	[-1, 48, 35, 35]	0
Conv2d-46	[-1, 64, 35, 35]	76,800
BatchNorm2d-47	[-1, 64, 35, 35]	128
BasicConv2d-48	[-1, 64, 35, 35]	0
Conv2d-49	[-1, 64, 35, 35]	16,384
BatchNorm2d-50	[-1, 64, 35, 35]	128
BasicConv2d-51	[-1, 64, 35, 35]	0
Conv2d-52	[-1, 96, 35, 35]	55,296

BatchNorm2d-53	[-1, 96, 35, 35]	192
BasicConv2d-54	[-1, 96, 35, 35]	0
Conv2d-55	[-1, 96, 35, 35]	82,944

```
data_transforms = {
    train: transforms.Compose([
        transforms.Resize((300,300)),
        transforms.RandomRotation(degrees=(0, 180)),
        transforms.RandomHorizontalFlip(0.5),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    test: transforms.Compose([
        transforms.Resize((300,300)),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
}
```

```
image_datasets = {x: datasets.ImageFolder(os.path.join(x), data_transforms[x]) for x in [t
```

```
dataloader = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=32, shuffle=True
```

```
dataset_sizes = {x: len(image_datasets[x]) for x in [train,test]}
```

```
class_names = image_datasets[train].classes
```

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554: UserWarning:
  warnings.warn(_create_warning_msg(
```

```
def train_batch(x, y, model, opt, loss_fn):
    output = model(x)
    # print(f"type of output - {type(output)}")
    batch_loss = loss_fn(output, y)
    batch_loss.backward()
    optimizer.step()
    optimizer.zero_grad()
    return batch_loss.item()
```

```
@torch.no_grad()
def accuracy(x, y, model):
    model.eval()
    prediction = model(x)
    max_values, argmaxes = prediction.max(-1)
    is_correct = argmaxes == y
    return is_correct.cpu().numpy().tolist()
```

```
@torch.no_grad()
def val_loss(x, y, model):
    model.eval()
    prediction = model(x)
    val_loss = loss_fn(prediction, y)
    return val_loss.item()

model, loss_fn, optimizer = get_model()

/usr/local/lib/python3.8/dist-packages/torchvision/models/_utils.py:208: UserWarning
  warnings.warn(
/usr/local/lib/python3.8/dist-packages/torchvision/models/_utils.py:223: UserWarning
  warnings.warn(msg)

train_losses, train_accuracies = [], []
val_losses, val_accuracies = [], []
for epoch in range(1,31):
    d0 = datetime.now()
    print(epoch)
    train_epoch_losses, train_epoch_accuracies = [], []
    for ix, batch in enumerate(iter(dataloader[train])):
#         print(f"ix - {ix}, {batch}")
        x, y = batch
#         print(f"type of x - {type(x)}, type of y - {type(y)}")
        x, y= x.to(device), y.to(device)
        batch_loss = train_batch(x, y, model, optimizer, loss_fn)
        is_correct = accuracy(x, y, model)
        train_epoch_accuracies.extend(is_correct)
        train_epoch_losses.append(batch_loss)
    train_epoch_loss = np.array(train_epoch_losses).mean()
    train_epoch_accuracy = np.mean(train_epoch_accuracies)
    print('Epoch:',epoch,'Train Loss:',train_epoch_loss,'Train Accuracy:',train_epoch_accu

    for ix, batch in enumerate(iter(dataloader[test])):
        x, y = batch
        x, y= x.to(device), y.to(device)
        val_is_correct = accuracy(x, y, model)
        validation_loss = val_loss(x, y, model)
        val_epoch_accuracy = np.mean(val_is_correct)
    dt = datetime.now() - d0
    print('Epoch:',epoch,'Validation Loss:',validation_loss,'Validation Accuracy:',val_epo


    train_losses.append(train_epoch_loss)
    train_accuracies.append(train_epoch_accuracy)
    val_losses.append(validation_loss)
    val_accuracies.append(val_epoch_accuracy)

1
Epoch: 1 Train Loss: 0.6147843960816877 Train Accuracy: 0.7997719823286304
```

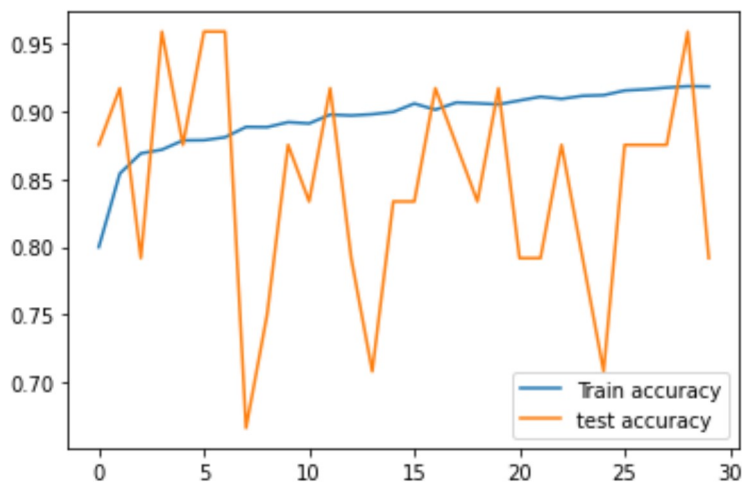

Epoch: 1 Validation Loss: 0.2536163628101349 Validation Accuracy: 0.875 Duration: 0:02
Epoch: 2 Train Loss: 0.4008501869263573 Train Accuracy: 0.8539261792788941
Epoch: 2 Validation Loss: 0.26119133830070496 Validation Accuracy: 0.9166666666666666
Epoch: 3 Train Loss: 0.3703767183260114 Train Accuracy: 0.8687473279179136
Epoch: 3 Validation Loss: 0.27270570397377014 Validation Accuracy: 0.7916666666666666
Epoch: 4 Train Loss: 0.35774958668306367 Train Accuracy: 0.8714550377654268
Epoch: 4 Validation Loss: 0.1237589493393898 Validation Accuracy: 0.9583333333333334
Epoch: 5 Train Loss: 0.3400826733856375 Train Accuracy: 0.8782955679065128
Epoch: 5 Validation Loss: 0.4670526683330536 Validation Accuracy: 0.875 Duration: 0:06
Epoch: 6 Train Loss: 0.33567620564362455 Train Accuracy: 0.8784380789511187
Epoch: 6 Validation Loss: 0.14849501848220825 Validation Accuracy: 0.9583333333333333
Epoch: 7 Train Loss: 0.3294972089672958 Train Accuracy: 0.880575744620208
Epoch: 7 Validation Loss: 0.18056625127792358 Validation Accuracy: 0.9583333333333333
Epoch: 8 Train Loss: 0.31252267434005043 Train Accuracy: 0.8882000855066268
Epoch: 8 Validation Loss: 1.1506531238555908 Validation Accuracy: 0.6666666666666666
Epoch: 9 Train Loss: 0.30893714871927924 Train Accuracy: 0.8880575744620208
Epoch: 9 Validation Loss: 0.6368729472160339 Validation Accuracy: 0.75 Duration: 0:10
Epoch: 10 Train Loss: 0.304772251932659 Train Accuracy: 0.8916916060994727
Epoch: 10 Validation Loss: 0.4377513825893402 Validation Accuracy: 0.875 Duration: 0:11
Epoch: 11 Train Loss: 0.3046975946966092 Train Accuracy: 0.890765284309534
Epoch: 11 Validation Loss: 0.772612988948822 Validation Accuracy: 0.8333333333333334
Epoch: 12 Train Loss: 0.2926434572523031 Train Accuracy: 0.897249536839105
Epoch: 12 Validation Loss: 0.22557474672794342 Validation Accuracy: 0.9166666666666666
Epoch: 13 Train Loss: 0.2871898828596216 Train Accuracy: 0.8966082371383782
Epoch: 13 Validation Loss: 0.7025075554847717 Validation Accuracy: 0.7916666666666666
Epoch: 14 Train Loss: 0.28246640831503617 Train Accuracy: 0.8975345589283169
Epoch: 14 Validation Loss: 0.5403130650520325 Validation Accuracy: 0.7083333333333333
Epoch: 15 Train Loss: 0.2787607992089148 Train Accuracy: 0.8992446914635884
Epoch: 15 Validation Loss: 0.6883854269981384 Validation Accuracy: 0.8333333333333333
Epoch: 16 Train Loss: 0.26896397121131826 Train Accuracy: 0.9053726663816446
Epoch: 16 Validation Loss: 0.5723026394844055 Validation Accuracy: 0.8333333333333333
Epoch: 17 Train Loss: 0.273323503463849 Train Accuracy: 0.9008835684765569
Epoch: 17 Validation Loss: 0.1693018078804016 Validation Accuracy: 0.9166666666666666
Epoch: 18 Train Loss: 0.2682374722678596 Train Accuracy: 0.9061564771269773
Epoch: 18 Validation Loss: 0.43931666016578674 Validation Accuracy: 0.875 Duration: 0:19
Epoch: 19 Train Loss: 0.2641515208913145 Train Accuracy: 0.9055864329485536
Epoch: 19 Validation Loss: 0.8991274833679199 Validation Accuracy: 0.8333333333333333

20

```
plt.plot(train_losses, label='train loss')
plt.plot(val_losses, label='test loss')
plt.legend()
plt.show()
```



```
plt.plot(train_accuracies, label='Train accuracy')
plt.plot(val_accuracies, label='test accuracy')
plt.legend()
plt.show()
```



```
n_correct = 0
n_total = 0
for inputs, targets in dataloader[train]:
    inputs, targets = inputs.to(device), targets.to(device)
    outputs = model(inputs)
    _, predictions = torch.max(outputs, -1)
    n_correct += (predictions == targets).sum().item()
    n_total += targets.shape[0]
train_acc = n_correct/n_total
```

```

n_correct = 0
n_total = 0
for inputs, targets in dataloader[test]:
    inputs, targets = inputs.to(device), targets.to(device)
    outputs = model(inputs)
    _, predictions = torch.max(outputs, -1)
    n_correct += (predictions == targets).sum().item()
    n_total += targets.shape[0]
test_acc = n_correct/n_total
print(f'Train acc: {train_acc:.4f}, Test acc: {test_acc:.4f}')

```

Train acc: 0.9111, Test acc: 0.8193

```

train_losses=[]
test_losses=[]
def accuracy(loader, model):
    num_corrects = 0
    num_samples = 0
    model.eval()
    loop = tqdm(loader)
    with torch.no_grad():
        for x, y in loop:
            x = x.to(device)
            y = y.to(device)
            scores = model(x)
            test_losses.append(scores.data)
            _, prediction = scores.max(1)
            num_corrects += (prediction == y).sum()
            num_samples += prediction.size(0)
            acc = (num_corrects/num_samples) * 100
            loop.set_postfix(acc=acc.item())
    print(f'Got {num_corrects}/{num_samples} with accuracy {acc:.4f}')

```

```
accuracy(dataloader[test], model)
```

100%

94/94 [00:10<00:00, 12.14it/s, acc=87]

Got 2609/3000 with accuracy 86.9667

Models summary

```
from torchsummary import summary
```

```

input_shape = (3,300,300)
summary(resnet.to(device), input_shape)

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 150, 150]	9,408
BatchNorm2d-2	[-1, 64, 150, 150]	128
ReLU-3	[-1, 64, 150, 150]	0
MaxPool2d-4	[-1, 64, 75, 75]	0
Conv2d-5	[-1, 64, 75, 75]	4,096
BatchNorm2d-6	[-1, 64, 75, 75]	128
ReLU-7	[-1, 64, 75, 75]	0
Conv2d-8	[-1, 64, 75, 75]	36,864
BatchNorm2d-9	[-1, 64, 75, 75]	128
ReLU-10	[-1, 64, 75, 75]	0
Conv2d-11	[-1, 256, 75, 75]	16,384
BatchNorm2d-12	[-1, 256, 75, 75]	512
Conv2d-13	[-1, 256, 75, 75]	16,384
BatchNorm2d-14	[-1, 256, 75, 75]	512
ReLU-15	[-1, 256, 75, 75]	0
Bottleneck-16	[-1, 256, 75, 75]	0
Conv2d-17	[-1, 64, 75, 75]	16,384
BatchNorm2d-18	[-1, 64, 75, 75]	128
ReLU-19	[-1, 64, 75, 75]	0
Conv2d-20	[-1, 64, 75, 75]	36,864
BatchNorm2d-21	[-1, 64, 75, 75]	128
ReLU-22	[-1, 64, 75, 75]	0
Conv2d-23	[-1, 256, 75, 75]	16,384
BatchNorm2d-24	[-1, 256, 75, 75]	512
ReLU-25	[-1, 256, 75, 75]	0
Bottleneck-26	[-1, 256, 75, 75]	0
Conv2d-27	[-1, 64, 75, 75]	16,384
BatchNorm2d-28	[-1, 64, 75, 75]	128
ReLU-29	[-1, 64, 75, 75]	0
Conv2d-30	[-1, 64, 75, 75]	36,864
BatchNorm2d-31	[-1, 64, 75, 75]	128
ReLU-32	[-1, 64, 75, 75]	0
Conv2d-33	[-1, 256, 75, 75]	16,384
BatchNorm2d-34	[-1, 256, 75, 75]	512
ReLU-35	[-1, 256, 75, 75]	0
Bottleneck-36	[-1, 256, 75, 75]	0
Conv2d-37	[-1, 128, 75, 75]	32,768
BatchNorm2d-38	[-1, 128, 75, 75]	256
ReLU-39	[-1, 128, 75, 75]	0
Conv2d-40	[-1, 128, 38, 38]	147,456
BatchNorm2d-41	[-1, 128, 38, 38]	256
ReLU-42	[-1, 128, 38, 38]	0
Conv2d-43	[-1, 512, 38, 38]	65,536
BatchNorm2d-44	[-1, 512, 38, 38]	1,024
Conv2d-45	[-1, 512, 38, 38]	131,072
BatchNorm2d-46	[-1, 512, 38, 38]	1,024
ReLU-47	[-1, 512, 38, 38]	0
Bottleneck-48	[-1, 512, 38, 38]	0
Conv2d-49	[-1, 128, 38, 38]	65,536
BatchNorm2d-50	[-1, 128, 38, 38]	256
ReLU-51	[-1, 128, 38, 38]	0
Conv2d-52	[-1, 128, 38, 38]	147,456

```

BatchNorm2d-53      [-1, 128, 38, 38]      256
ReLU-54             [-1, 128, 38, 38]      0
Conv2d-55           [-1, 512, 38, 38]      65,536

```

```
from torchsummary import summary
```

```
input_shape = (3,300,300)
summary(wrn.to(device), input_shape)
```

```

-----
Layer (type)          Output Shape          Param #
-----
Conv2d-1              [-1, 64, 150, 150]    9,408
BatchNorm2d-2         [-1, 64, 150, 150]    128
ReLU-3                [-1, 64, 150, 150]    0
MaxPool2d-4           [-1, 64, 75, 75]      0
Conv2d-5              [-1, 128, 75, 75]     8,192
BatchNorm2d-6         [-1, 128, 75, 75]     256
ReLU-7               [-1, 128, 75, 75]     0
Conv2d-8              [-1, 128, 75, 75]    147,456
BatchNorm2d-9         [-1, 128, 75, 75]     256
ReLU-10              [-1, 128, 75, 75]     0
Conv2d-11             [-1, 256, 75, 75]    32,768
BatchNorm2d-12        [-1, 256, 75, 75]     512
Conv2d-13             [-1, 256, 75, 75]    16,384
BatchNorm2d-14        [-1, 256, 75, 75]     512
ReLU-15              [-1, 256, 75, 75]     0
Bottleneck-16        [-1, 256, 75, 75]     0
Conv2d-17             [-1, 128, 75, 75]    32,768
BatchNorm2d-18        [-1, 128, 75, 75]     256
ReLU-19              [-1, 128, 75, 75]     0
Conv2d-20             [-1, 128, 75, 75]    147,456
BatchNorm2d-21        [-1, 128, 75, 75]     256
ReLU-22              [-1, 128, 75, 75]     0
Conv2d-23             [-1, 256, 75, 75]    32,768
BatchNorm2d-24        [-1, 256, 75, 75]     512
ReLU-25              [-1, 256, 75, 75]     0
Bottleneck-26        [-1, 256, 75, 75]     0
Conv2d-27             [-1, 128, 75, 75]    32,768
BatchNorm2d-28        [-1, 128, 75, 75]     256
ReLU-29              [-1, 128, 75, 75]     0
Conv2d-30             [-1, 128, 75, 75]    147,456
BatchNorm2d-31        [-1, 128, 75, 75]     256
ReLU-32              [-1, 128, 75, 75]     0
Conv2d-33             [-1, 256, 75, 75]    32,768
BatchNorm2d-34        [-1, 256, 75, 75]     512
ReLU-35              [-1, 256, 75, 75]     0
Bottleneck-36        [-1, 256, 75, 75]     0
Conv2d-37             [-1, 256, 75, 75]    65,536
BatchNorm2d-38        [-1, 256, 75, 75]     512
ReLU-39              [-1, 256, 75, 75]     0
Conv2d-40             [-1, 256, 38, 38]    589,824
BatchNorm2d-41        [-1, 256, 38, 38]     512
ReLU-42              [-1, 256, 38, 38]     0
Conv2d-43             [-1, 512, 38, 38]    121,072

```

Conv2d-43	[-1, 512, 38, 38]	131,072
BatchNorm2d-44	[-1, 512, 38, 38]	1,024
Conv2d-45	[-1, 512, 38, 38]	131,072
BatchNorm2d-46	[-1, 512, 38, 38]	1,024
ReLU-47	[-1, 512, 38, 38]	0
Bottleneck-48	[-1, 512, 38, 38]	0
Conv2d-49	[-1, 256, 38, 38]	131,072
BatchNorm2d-50	[-1, 256, 38, 38]	512
ReLU-51	[-1, 256, 38, 38]	0
Conv2d-52	[-1, 256, 38, 38]	589,824
BatchNorm2d-53	[-1, 256, 38, 38]	512
ReLU-54	[-1, 256, 38, 38]	0
Conv2d-55	[-1, 512, 38, 38]	131,072

```
from torchsummary import summary
```

```
input_shape = (3,300,300)
```

```
summary(model.to(device), input_shape)
```

```
-----
```

Layer (type)	Output Shape	Param #
-----	-----	-----
Conv2d-1	[-1, 32, 149, 149]	864
BatchNorm2d-2	[-1, 32, 149, 149]	64
BasicConv2d-3	[-1, 32, 149, 149]	0
Conv2d-4	[-1, 32, 147, 147]	9,216
BatchNorm2d-5	[-1, 32, 147, 147]	64
BasicConv2d-6	[-1, 32, 147, 147]	0
Conv2d-7	[-1, 64, 147, 147]	18,432
BatchNorm2d-8	[-1, 64, 147, 147]	128
BasicConv2d-9	[-1, 64, 147, 147]	0
MaxPool2d-10	[-1, 64, 73, 73]	0
Conv2d-11	[-1, 80, 73, 73]	5,120
BatchNorm2d-12	[-1, 80, 73, 73]	160
BasicConv2d-13	[-1, 80, 73, 73]	0
Conv2d-14	[-1, 192, 71, 71]	138,240
BatchNorm2d-15	[-1, 192, 71, 71]	384
BasicConv2d-16	[-1, 192, 71, 71]	0
MaxPool2d-17	[-1, 192, 35, 35]	0
Conv2d-18	[-1, 64, 35, 35]	12,288
BatchNorm2d-19	[-1, 64, 35, 35]	128
BasicConv2d-20	[-1, 64, 35, 35]	0
Conv2d-21	[-1, 48, 35, 35]	9,216
BatchNorm2d-22	[-1, 48, 35, 35]	96
BasicConv2d-23	[-1, 48, 35, 35]	0
Conv2d-24	[-1, 64, 35, 35]	76,800
BatchNorm2d-25	[-1, 64, 35, 35]	128
BasicConv2d-26	[-1, 64, 35, 35]	0
Conv2d-27	[-1, 64, 35, 35]	12,288
BatchNorm2d-28	[-1, 64, 35, 35]	128
BasicConv2d-29	[-1, 64, 35, 35]	0
Conv2d-30	[-1, 96, 35, 35]	55,296
BatchNorm2d-31	[-1, 96, 35, 35]	192
BasicConv2d-32	[-1, 96, 35, 35]	0
Conv2d-33	[-1, 96, 35, 35]	82,944

BatchNorm2d-34	[-1, 96, 35, 35]	192
BasicConv2d-35	[-1, 96, 35, 35]	0
Conv2d-36	[-1, 32, 35, 35]	6,144
BatchNorm2d-37	[-1, 32, 35, 35]	64
BasicConv2d-38	[-1, 32, 35, 35]	0
InceptionA-39	[-1, 256, 35, 35]	0
Conv2d-40	[-1, 64, 35, 35]	16,384
BatchNorm2d-41	[-1, 64, 35, 35]	128
BasicConv2d-42	[-1, 64, 35, 35]	0
Conv2d-43	[-1, 48, 35, 35]	12,288
BatchNorm2d-44	[-1, 48, 35, 35]	96
BasicConv2d-45	[-1, 48, 35, 35]	0
Conv2d-46	[-1, 64, 35, 35]	76,800
BatchNorm2d-47	[-1, 64, 35, 35]	128
BasicConv2d-48	[-1, 64, 35, 35]	0
Conv2d-49	[-1, 64, 35, 35]	16,384
BatchNorm2d-50	[-1, 64, 35, 35]	128
BasicConv2d-51	[-1, 64, 35, 35]	0
Conv2d-52	[-1, 96, 35, 35]	55,296
BatchNorm2d-53	[-1, 96, 35, 35]	192
BasicConv2d-54	[-1, 96, 35, 35]	0
Conv2d-55	[-1, 96, 35, 35]	82,944

[Colab paid products](#) - [Cancel contracts here](#)